

Advanced Machine Learning Techniques

Gradient Descent, Errors

Accessing the dataset

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: # Access the olympics dataset

oly=pd.read_csv('olympics.csv')
```

```
In [3]: oly
```

```
Out[3]:
```

	Year	Time
0	1896	12.00
1	1900	11.00
2	1904	11.00
3	1908	10.80
4	1912	10.80
5	1920	10.80
6	1924	10.60
7	1928	10.80
8	1932	10.38
9	1936	10.30
10	1948	10.30
11	1952	10.40
12	1956	10.50
13	1960	10.20
14	1964	10.00
15	1968	9.95
16	1972	10.14
17	1976	10.06
18	1980	10.25
19	1984	9.99
20	1988	9.92
21	1992	9.96
22	1996	9.84
23	2000	9.87
24	2004	9.85
25	2008	9.69
26	2012	9.63
27	2016	9.81
28	2021	9.80

Preprocessing

```
In [4]: X=oly['Year']  
X
```


```
Out[4]: 0      1896  
1      1900  
2      1904  
3      1908  
4      1912  
5      1920  
6      1924  
7      1928  
8      1932  
9      1936  
10     1948  
11     1952  
12     1956  
13     1960  
14     1964  
15     1968  
16     1972  
17     1976  
18     1980  
19     1984  
20     1988  
21     1992  
22     1996  
23     2000  
24     2004  
25     2008  
26     2012  
27     2016  
28     2021  
Name: Year, dtype: int64
```

```
In [5]: y=oly['Time']  
y
```

```
Out[5]: 0      12.00  
1      11.00  
2      11.00  
3      10.80  
4      10.80  
5      10.80  
6      10.60  
7      10.80  
8      10.38  
9      10.30  
10     10.30  
11     10.40  
12     10.50  
13     10.20  
14     10.00  
15       9.95  
16     10.14  
17     10.06  
18     10.25  
19       9.99  
20       9.92  
21       9.96  
22       9.84  
23       9.87  
24       9.85  
25       9.69  
26       9.63  
27       9.81  
28       9.80  
Name: Time, dtype: float64
```

Plotting

```
In [7]: plt.figure(figsize=(10,8))  
plt.scatter(X,y);
```




Desinging a simple LR model


```
In [10]: from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()  
  
lr.fit(X.to_numpy().reshape(-1,1),y)
```

```
Out[10]: LinearRegression()
```

```
In [11]: # Plotting  
  
plt.figure(figsize=(10,8))  
plt.scatter(X,y)  
plt.plot(X,lr.predict(X.to_numpy().reshape(-1,1)));
```



What is behind?



Error:

Bias:

Variance

The model underfits

Polynomial Regression

Degree 2

```
In [12]: from sklearn.preprocessing import PolynomialFeatures

p2=PolynomialFeatures(degree=2)

Xp2=p2.fit_transform(X.to_numpy().reshape(-1,1))

Xp2
```

```
Out[12]: array([[1.000000e+00, 1.896000e+03, 3.594816e+06],
 [1.000000e+00, 1.900000e+03, 3.610000e+06],
 [1.000000e+00, 1.904000e+03, 3.625216e+06],
 [1.000000e+00, 1.908000e+03, 3.640464e+06],
 [1.000000e+00, 1.912000e+03, 3.655744e+06],
 [1.000000e+00, 1.920000e+03, 3.686400e+06],
 [1.000000e+00, 1.924000e+03, 3.701776e+06],
 [1.000000e+00, 1.928000e+03, 3.717184e+06],
 [1.000000e+00, 1.932000e+03, 3.732624e+06],
 [1.000000e+00, 1.936000e+03, 3.748096e+06],
 [1.000000e+00, 1.948000e+03, 3.794704e+06],
 [1.000000e+00, 1.952000e+03, 3.810304e+06],
 [1.000000e+00, 1.956000e+03, 3.825936e+06],
 [1.000000e+00, 1.960000e+03, 3.841600e+06],
 [1.000000e+00, 1.964000e+03, 3.857296e+06],
 [1.000000e+00, 1.968000e+03, 3.873024e+06],
 [1.000000e+00, 1.972000e+03, 3.888784e+06],
 [1.000000e+00, 1.976000e+03, 3.904576e+06],
 [1.000000e+00, 1.980000e+03, 3.920400e+06],
 [1.000000e+00, 1.984000e+03, 3.936256e+06],
 [1.000000e+00, 1.988000e+03, 3.952144e+06],
 [1.000000e+00, 1.992000e+03, 3.968064e+06],
 [1.000000e+00, 1.996000e+03, 3.984016e+06],
 [1.000000e+00, 2.000000e+03, 4.000000e+06],
 [1.000000e+00, 2.004000e+03, 4.016016e+06],
 [1.000000e+00, 2.008000e+03, 4.032064e+06],
 [1.000000e+00, 2.012000e+03, 4.048144e+06],
 [1.000000e+00, 2.016000e+03, 4.064256e+06],
 [1.000000e+00, 2.021000e+03, 4.084441e+06]])
```

In [13]:


X

```
Out[13]: 0    1896
          1    1900
          2    1904
          3    1908
          4    1912
          5    1920
          6    1924
          7    1928
          8    1932
          9    1936
         10    1948
         11    1952
         12    1956
         13    1960
         14    1964
         15    1968
         16    1972
         17    1976
         18    1980
         19    1984
         20    1988
         21    1992
         22    1996
         23    2000
         24    2004
         25    2008
         26    2012
         27    2016
         28    2021
          Name: Year, dtype: int64
```

```
In [14]: lr2=LinearRegression()
          pr2=lr2.fit(Xp2,y)
```

In [17]: # Plotting

```
plt.figure(figsize=(10,8))
plt.scatter(X,y)
plt.plot(X,lr.predict(X.to_numpy().reshape(-1,1)),label='Linear Reg')
plt.plot(X,pr2.predict(Xp2), label='Poly Reg degree 2' )
plt.legend();
```



```
In [19]: # Degree 3

p3=PolynomialFeatures(degree=3)


Xp3=p3.fit_transform(X.to_numpy().reshape(-1,1))

Xp3

lr3= LinearRegression()
pr3=lr3.fit(Xp3,y)
```

```
In [21]: # Plotting

plt.figure(figsize=(10,8))
plt.scatter(X,y)
plt.plot(X,lr.predict(X.to_numpy().reshape(-1,1)),label='Linear Reg')
plt.plot(X,pr2.predict(Xp2), label='Poly Reg degree 2' )
plt.plot(X,pr3.predict(Xp3), label='Poly Reg degree 3')
plt.legend();
```



```
In [22]: # Degree 5


p5=PolynomialFeatures(degree=5)

Xp5=p5.fit_transform(X.to_numpy().reshape(-1,1))

lr5= LinearRegression()
pr5=lr5.fit(Xp5,y)
```

In [23]: # Plotting

```
plt.figure(figsize=(10,8))
plt.scatter(X,y)
plt.plot(X,lr.predict(X.to_numpy().reshape(-1,1)),label='Linear Reg')
plt.plot(X,pr2.predict(Xp2), label='Poly Reg degree 2' )
plt.plot(X,pr3.predict(Xp3), label='Poly Reg degree 3')
plt.plot(X,pr5.predict(Xp5), label='Poly Reg degree 5')
plt.legend();
```




In [24]: # Degree 10

```
p10=PolynomialFeatures(degree=10)
Xp10=p10.fit_transform(X.to_numpy().reshape(-1,1))

lr10= LinearRegression()
pr10=lr10.fit(Xp10,y)
```

In [25]: # Plotting

```
plt.figure(figsize=(10,8))
plt.scatter(X,y)
plt.plot(X,lr.predict(X.to_numpy().reshape(-1,1)),label='Linear Reg')
plt.plot(X,pr2.predict(Xp2), label='Poly Reg degree 2' )
plt.plot(X,pr3.predict(Xp3), label='Poly Reg degree 3')
plt.plot(X,pr5.predict(Xp5), label='Poly Reg degree 5')
plt.plot(X,pr10.predict(Xp10), label='Poly Reg degree 10')
plt.legend();
```

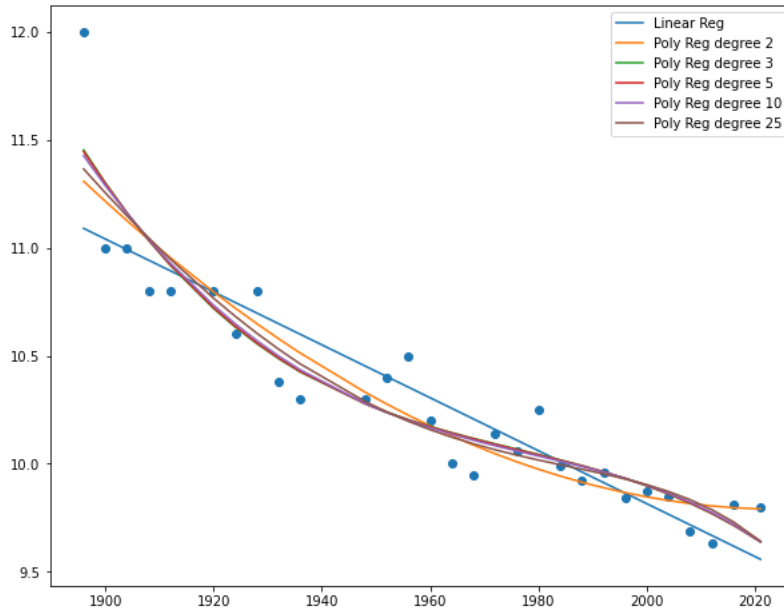


In [26]: # Degree 25

```
p25=PolynomialFeatures(degree=25)
Xp25=p25.fit_transform(X.to_numpy().reshape(-1,1))
lr25= LinearRegression()
pr25=lr25.fit(Xp25,y)
```


In [27]: # Plotting


```
plt.figure(figsize=(10,8))
plt.scatter(X,y)
plt.plot(X,lr.predict(X.to_numpy().reshape(-1,1)),label='Linear Reg')
plt.plot(X,pr2.predict(Xp2), label='Poly Reg degree 2' )
plt.plot(X,pr3.predict(Xp3), label='Poly Reg degree 3')
plt.plot(X,pr5.predict(Xp5), label='Poly Reg degree 5')
plt.plot(X,pr10.predict(Xp10), label='Poly Reg degree 10')
plt.plot(X,pr25.predict(Xp25), label='Poly Reg degree 25')
plt.legend();
```




In [28]: # Degree 75

```
p75=PolynomialFeatures(degree=75)
Xp75=p75.fit_transform(X.to_numpy().reshape(-1,1))
lr75= LinearRegression()
pr75=lr75.fit(Xp75,y)
```

```
In [29]: plt.figure(figsize=(10,8))
plt.scatter(X,y)
plt.plot(X,lr.predict(X.to_numpy().reshape(-1,1)),label='Linear Reg')
plt.plot(X,pr2.predict(Xp2), label='Poly Reg degree 2' )
plt.plot(X,pr3.predict(Xp3), label='Poly Reg degree 3')
plt.plot(X,pr5.predict(Xp5), label='Poly Reg degree 5')
plt.plot(X,pr10.predict(Xp10), label='Poly Reg degree 10')
plt.plot(X,pr25.predict(Xp25), label='Poly Reg degree 25')
plt.plot(X,pr75.predict(Xp75), label='Poly Reg degree 75')
plt.legend();
```



```
In [31]: plt.figure(figsize=(10,8))
plt.scatter(X,y)
plt.plot(X,lr.predict(X.to_numpy().reshape(-1,1)),label='Linear Reg')
# plt.plot(X,pr2.predict(Xp2), label='Poly Reg degree 2' )
# plt.plot(X,pr3.predict(Xp3), label='Poly Reg degree 3')
# plt.plot(X,pr5.predict(Xp5), label='Poly Reg degree 5')
# plt.plot(X,pr10.predict(Xp10), label='Poly Reg degree 10')
# plt.plot(X,pr25.predict(Xp25), label='Poly Reg degree 25')
plt.plot(X,pr75.predict(Xp75), label='Poly Reg degree 75')
plt.legend();
```



Errors:


Bias error

Corresponding the model selected >>> results in underfitting

Variance error

Corresponding to the training data >>> results overfitting

Bias variance trade-off



In []: